

ARUL SELVAM P

New-Manuscript-CPE-25-0596-Rev1-5_R5fixes_Clean1.docx

 My Files My Files University

Document Details

Submission ID

trn:oid:::6447:308878865

Submission Date

Sep 7, 2025, 3:45 PM GMT+5:30

Download Date

Sep 7, 2025, 3:47 PM GMT+5:30

File Name

New-Manuscript-CPE-25-0596-Rev1-5_R5fixes_Clean1.docx

File Size

683.6 KB

21 Pages

6,909 Words

45,561 Characters





9% Overall Similarity

The combined total of all matches, including overlapping sources, for each database.




Filtered from the Report

- Bibliography

Match Groups

-  **45 Not Cited or Quoted** 6%
Matches with neither in-text citation nor quotation marks
-  **14 Missing Quotations** 2%
Matches that are still very similar to source material
-  **0 Missing Citation** 0%
Matches that have quotation marks, but no in-text citation
-  **0 Cited and Quoted** 0%
Matches with in-text citation present, but no quotation marks

Top Sources

- 4%  Internet sources
- 4%  Publications
- 6%  Submitted works (Student Papers)

Integrity Flags

0 Integrity Flags for Review

No suspicious text manipulations found.

Our system's algorithms look deeply at a document for any inconsistencies that would set it apart from a normal submission. If we notice something strange, we flag it for you to review.

A Flag is not necessarily an indicator of a problem. However, we'd recommend you focus your attention there for further review.

Match Groups

- 45 Not Cited or Quoted** 6%
Matches with neither in-text citation nor quotation marks
- 14 Missing Quotations** 2%
Matches that are still very similar to source material
- 0 Missing Citation** 0%
Matches that have quotation marks, but no in-text citation
- 0 Cited and Quoted** 0%
Matches with in-text citation present, but no quotation marks

Top Sources

- 4% Internet sources
- 4% Publications
- 6% Submitted works (Student Papers)

Top Sources

The sources with the highest number of matches within the submission. Overlapping sources will not be displayed.

1	Internet	onlinelibrary.wiley.com	<1%
2	Internet	www.mdpi.com	<1%
3	Internet	aimlstudies.co.uk	<1%
4	Internet	ajates-scholarly.com	<1%
5	Submitted works	University of Lagos on 2025-08-27	<1%
6	Submitted works	Dublin Business School on 2025-08-28	<1%
7	Submitted works	University of Wales, Bangor on 2024-05-30	<1%
8	Internet	sos-vo.org	<1%
9	Publication	Ramesh Babu Chellappan. "From Algorithms to Accountability: The Societal and E...	<1%
10	Internet	notionpress.com	<1%

11	Submitted works	University of Warwick on 2025-09-01	<1%
12	Publication	Manyar, Omey M.. "Physics-Informed AI Methods for Deformable Object Manipul...	<1%
13	Publication	Birudala Venkatesh Reddy, Y V Krishna Reddy, Md. Abdur Razzak, Surender Reddy...	<1%
14	Submitted works	Brickfields Asia College on 2025-01-15	<1%
15	Submitted works	ESoft Metro Campus, Sri Lanka on 2025-08-24	<1%
16	Submitted works	Swiss School of Business and Management - SSBM on 2025-07-17	<1%
17	Submitted works	Texas A&M University, College Station on 2023-05-02	<1%
18	Internet	www2.mdpi.com	<1%
19	Submitted works	Rochester Institute of Technology on 2025-04-27	<1%
20	Publication	Xinyuan Song, HSIEH,WEI-CHE, Ziqian Bi, Chuanqi Jiang, Junyu Liu, Benji Peng, Se...	<1%
21	Internet	hal.science	<1%
22	Internet	professions.ng	<1%
23	Publication	Krishan Kumar, Kiran Jyoti. "Enhancing Transparency and Trust in Brain Tumor Di...	<1%
24	Submitted works	Tilburg University on 2025-06-05	<1%

25	Internet	arxiv.org	<1%
26	Internet	ijircce.com	<1%
27	Internet	ijitce.org	<1%
28	Internet	www.fastercapital.com	<1%
29	Internet	www.medrxiv.org	<1%
30	Submitted works	Dublin Business School on 2025-08-27	<1%
31	Publication	Ennab, Mohammad. "A Hybrid Convolutional-Fuzzy Model for Interpretable AI in ...	<1%
32	Publication	Janner, Michael. "Deep Generative Models for Decision-Making and Control", Univ...	<1%
33	Publication	Mourade Azrour, Jamal Mabrouki, Sultan Ahmad. "IoT and Advanced Intelligence ...	<1%
34	Publication	Pawan Singh Mehra, Dharendra Kumar Shukla. "Artificial Intelligence, Blockchain,...	<1%
35	Submitted works	Sheffield Hallam University on 2025-09-04	<1%
36	Submitted works	University of Bradford on 2025-09-02	<1%
37	Submitted works	University of Essex on 2025-09-03	<1%
38	Publication	Yogesh K. Dwivedi, Nir Kshetri, Laurie Hughes, Emma Louise Slade et al. ""So wha...	<1%

39	Internet	assets-eu.researchsquare.com	<1%
40	Internet	data-science.semuniver.kz	<1%
41	Internet	ebin.pub	<1%
42	Internet	lucris.lub.lu.se	<1%
43	Internet	openreviewhub.org	<1%
44	Internet	www.frontiersin.org	<1%
45	Submitted works	Liverpool John Moores University on 2025-05-25	<1%
46	Publication	S. Kannadhasan, R. Nagarajan, Alagar Karthick, V. Kumar Chinnaiyan. "Technolog...	<1%
47	Submitted works	Swiss School of Business and Management - SSBM on 2025-08-23	<1%
48	Submitted works	BBPlugin on 2025-08-22	<1%
49	Publication	John R. Vacca. "Cloud Computing Security - Foundations and Challenges", CRC Pre...	<1%
50	Publication	Pushpa Choudhary, Sambit Satpathy, Arvind Dagur, Dharendra Kumar Shukla. "Re...	<1%
51	Submitted works	University of Derby on 2014-12-13	<1%

Explainable AI (XAI) for Insider Threat Detection: Balancing Security and Transparency in Cloud Computing

ARUL SELVAM P^{1,*} and TAMIJE SELVY P²

¹ *Department of Artificial Intelligence and Machine Learning, Hindusthan College of Engineering and Technology, Coimbatore, Tamil Nadu, India.*

² *Department of Computer Science and Engineering, Hindusthan College of Engineering and Technology, Coimbatore, Tamil Nadu, India.*

*Correspondence: Arul Selvam P, *Department of Artificial Intelligence and Machine Learning, Hindusthan College of Engineering and Technology, Coimbatore, Tamil Nadu, India. Email: arulselvamme@gmail.com*

Abstract

Mitigating insider threats in cloud environments remains highly challenging, as malicious or negligent activities frequently evade traditional defense mechanisms. Although AI-based models can detect such behaviors effectively, their lack of interpretability often reduces analyst confidence and hinders adoption in mission-critical settings. To overcome this limitation, we introduce a hybrid detection framework that combines anomaly detection with explainable artificial intelligence (XAI). The framework integrates Shapley Additive Explanations (SHAP), Local Interpretable Model-Agnostic Explanations (LIME), and counterfactual reasoning to provide transparent and actionable insights into model decisions. Furthermore, we define two quantitative measures—Explainability Score (ES) and Fidelity Score (FS)—to systematically assess interpretability alongside predictive accuracy. Experimental validation on benchmark datasets (CERT, ADFA-LD) and large-scale synthetic logs demonstrates that the proposed system achieves $94.5\% \pm 0.4$ accuracy, reduces false positives by 27%, and significantly enhances analyst efficiency. In contrast to conventional black-box models, the framework strikes a practical balance between accuracy and interpretability, thereby improving trustworthiness and reinforcing security operations in the cloud.

Keywords: *Explainable AI (XAI), Insider Threat Detection, Cloud Security, SHAP, LIME, Counterfactual Reasoning, Anomaly Detection*

1. Introduction

Cloud computing has rapidly evolved into a foundational technology for enterprises, offering scalability, efficiency, and improved service delivery. Yet, its openness and reliance on distributed access models expose organizations to substantial cybersecurity challenges, particularly insider threats. These arise when employees, contractors, or partners abuse their legitimate access to compromise data, systems, or services [8]. Unlike external attacks, insider incidents are harder to detect since they exploit valid credentials, masking malicious activity as routine operations [10].

Artificial intelligence (AI) and machine learning (ML) have gained traction for detecting insider threats by learning behavioral baselines, identifying anomalies, and automating responses [11]. However, a persistent obstacle is that most AI models behave as opaque “black boxes” [9]. Their lack of interpretability undermines analyst trust, complicates compliance

obligations, and reduces confidence in automated alerts [4]. Explainable AI (XAI) offers a solution by combining predictive power with interpretability. Techniques such as Shapley Additive Explanations (SHAP), Local Interpretable Model-Agnostic Explanations (LIME), and counterfactual reasoning expose the rationale behind model predictions, enabling analysts to validate alerts and improve decision-making [1]. Unlike earlier studies that applied these methods in isolation, this work unifies SHAP, LIME, and counterfactual reasoning into a single insider threat detection framework that balances accuracy with transparency.

1.1 Background and Motivation

The benefits of cloud adoption are accompanied by growing vulnerabilities. According to a recent Cloud Security Alliance (CSA) survey, nearly 60% of cloud-related breaches originate from insider activity [3]. Traditional detection methods—such as rule-based or signature-driven systems—struggle to keep pace with evolving insider behaviors that easily evade static defenses [7].

As a result, AI-based approaches have gained importance. By analyzing authentication events, access logs, and user behavior patterns, ML models can identify deviations that suggest insider misuse [12]. Still, most of these approaches remain opaque, preventing analysts from understanding the basis of their alerts [6]. This interpretability gap reduces trust in AI-driven security decisions and creates difficulties in complying with regulations such as the General Data Protection Regulation (GDPR) [5] and the National Institute of Standards and Technology's AI Risk Management Framework (NIST AI RMF 1.0, 2023) [6]. Both emphasize transparency and accountability in automated decision-making. XAI addresses this shortfall by clarifying model logic, thereby enabling both operational validation and regulatory alignment [2].

1.2 Research Problem Statement

Insider threat detection faces a fundamental tension: while deep learning and other advanced AI models achieve strong accuracy, their decision processes remain opaque—the so-called “black-box problem” [4]. This challenge raises three main issues:

1. **Trust and Explainability** – Security analysts hesitate to act on opaque alerts without insight into the underlying rationale.
2. **Regulatory Compliance** – Data protection frameworks increasingly require transparency in automated decision-making.
3. **Operational Effectiveness** – Without interpretability, organizations cannot easily assess risks or formulate targeted responses.

Although anomaly detection methods have advanced, most insider threat detection systems in cloud contexts still fail to bridge the gap between accuracy and interpretability. This disconnect motivates the present research: to design a detection framework that provides both high predictive performance and actionable transparency.

1.3 Research Objectives and Contributions

To address these challenges, this study sets out the following objectives:

- **Framework Development** – Design an XAI-based insider threat detection framework that enhances both transparency and security.
- **Integration of Interpretable Techniques** – Employ SHAP, LIME, and counterfactual reasoning to generate clear, analyst-friendly explanations.
- **Comprehensive Evaluation** – Validate the framework on multiple datasets, including CERT, ADFA-LD, and large-scale synthetic cloud logs.

The major contributions are:

- **Hybrid XAI Framework** – A unified insider threat detection model combining anomaly detection with SHAP, LIME, and counterfactual reasoning.
- **Explainability Metrics** – Introduction of two measures—Explainability Score (ES) and Fidelity Score (FS)—to quantify interpretability alongside accuracy.
- **Comparative Evaluation** – Benchmarking against black-box and XAI-enhanced baselines, showing balanced trade-offs between detection accuracy and interpretability (see Tables 1–2).
- **Analyst Utility** – Demonstration that layered XAI explanations reduce false positives by 27% and shorten analyst investigation time, thereby improving operational decision-making.

By advancing explainable detection, this research provides a practical, trustworthy approach to mitigating insider threats in enterprise cloud environments.

2. Literature Review

The widespread shift to cloud computing has reshaped how organizations store, process, and share information. While it provides scalability and flexibility, it also amplifies security concerns—particularly insider threats, which are typically harder to detect than external intrusions. This section reviews prior research on insider threats in cloud environments, AI-based detection approaches, the contribution of explainable AI (XAI) to improving transparency, and remaining challenges that guide this study.

2.1 Insider Threats in Cloud Computing

Insider threats occur when individuals with authorized access misuse their privileges to compromise organizational assets [8]. Unlike outside attackers, insiders operate within trusted boundaries using legitimate credentials, making their activity resemble normal behavior and difficult to flag [10].

Scholars generally classify insiders into three groups: malicious insiders, who intentionally cause harm through theft or sabotage; negligent insiders, whose mistakes (e.g., weak authentication practices, misconfigurations) introduce vulnerabilities; and compromised insiders, whose accounts are hijacked and exploited by external actors [7].

These threats manifest in various ways, such as privilege escalation leading to unauthorized data exfiltration [12], inadvertent leaks due to poor access management, or account compromise enabling adversaries to blend into standard workflows [9]. Because such incidents imitate legitimate user behavior, conventional monitoring tools often fail to distinguish them, prompting the need for advanced, behavior-aware approaches.

2.2 AI-Based Insider Threat Detection

Artificial Intelligence (AI) and Machine Learning (ML) techniques have emerged as central to detecting insider activity in cloud systems. By mining logs, authentication trails, and access records, AI models learn behavioral baselines and flag deviations indicative of threats [5].

Detection typically follows two paradigms. Supervised learning uses labeled datasets to train classifiers that differentiate benign from malicious behavior. While effective when labeled data is available, this approach is constrained by the scarcity of annotated insider threat examples [6]. Unsupervised learning, by contrast, does not depend on labels and instead identifies anomalies in activity. Methods such as clustering, autoencoders, and isolation forests detect patterns like off-hours logins, unusual data transfers, or abnormal privilege use [4].

Despite successes, limitations remain. Many ML-driven systems still function as black boxes, offering little insight into why a prediction was made [1]. Data imbalance can produce biased models, undermining fairness and robustness [11]. Moreover, high false-positive rates continue to overwhelm analysts, lowering practical adoption [7]. These issues highlight the need for solutions that combine predictive strength with interpretability.

2.3 Explainable AI (XAI) in Security

XAI provides mechanisms to clarify how AI models generate outputs, a feature especially important in cloud security where analyst trust and regulatory compliance are paramount [2].

Different XAI approaches have been applied in this context. SHAP (Shapley Additive Explanations) attributes importance values to input features, highlighting key factors such as unusual login hours or large file transfers [9]. LIME (Local Interpretable Model-Agnostic Explanations) constructs simplified local models that explain individual predictions, helping analysts validate why specific alerts—such as an unexpected login combined with data movement—were triggered [6]. Counterfactual reasoning introduces “what-if” analysis, showing minimal behavioral changes that would alter a classification outcome, aiding in differentiating harmless anomalies from actual threats [4].

Symbolic systems (e.g., rule-based engines) offer interpretability but lack flexibility, while deep neural networks achieve high accuracy yet remain opaque [12]. XAI bridges this gap by combining predictive accuracy with explanations that meet compliance requirements and enhance analyst usability.

2.4 Gaps in Existing Research

Although advances have been made, several gaps remain in the application of XAI to insider threat detection. Scalability is a key concern: many explanation methods work in controlled experiments but struggle in distributed, real-time cloud environments [11]. Privacy is another challenge, since explanation outputs can inadvertently expose sensitive identifiers or activity details. Future frameworks must therefore incorporate privacy-preserving methods to align with standards such as GDPR [5].

Another limitation is the reliance on synthetic or laboratory datasets. Many studies lack validation in operational environments, leaving uncertainties about their deployment readiness

[7]. Addressing these gaps requires solutions that are both computationally efficient and empirically validated in real-world cloud infrastructures.

3. Proposed XAI-Based Insider Threat Detection Framework

Detecting insider threats remains a persistent difficulty in cloud environments, as such risks arise from legitimate accounts misused for malicious or negligent purposes. Traditional AI-based anomaly detectors can identify suspicious behavior but are often deployed as opaque black-box models, hindering analyst confidence and slowing incident response. To address this, we present an explainable AI (XAI)-based framework that integrates behavior monitoring, anomaly detection, and multi-level interpretability. The objective is to deliver security alerts that are both precise and transparent, allowing analysts to make timely and well-informed decisions.

3.1 System Architecture

The proposed system is organized in three functional layers (Figure 1), each designed to improve detection capability while ensuring interpretability of results.

- **User Behavior Analytics (UBA):** This layer consolidates diverse logs such as authentication attempts, file operations, process activity, and network connections. Individual user baselines are established, and deviations from these profiles are treated as candidate threats.
- **Anomaly Detection with Interpretability:** Machine learning (ML) and deep learning (DL) techniques are used to uncover deviations from expected behavior. To reduce the opacity of such models, the framework integrates Shapley Additive Explanations (SHAP) and Local Interpretable Model-Agnostic Explanations (LIME) [20,25], enabling analysts to see which features most influenced a given alert and to differentiate false alarms from genuine threats.
- **Integration with Enterprise Security Tools:** The pipeline is designed to operate within existing enterprise security infrastructures:
 - *SIEM (Security Information and Event Management)* for centralized event collection and correlation,
 - *IAM (Identity and Access Management)* for adaptive access control, and
 - *CASB (Cloud Access Security Broker)* for policy enforcement at the cloud service level [11].

By combining these layers into a cohesive workflow, the framework ensures that insider threat detection achieves both strong predictive performance and operational interpretability.

3.2 XAI Integration for Threat Interpretation

To make alerts actionable, the framework embeds three complementary XAI techniques (Figure 2), each offering a distinct interpretive perspective:

- **SHAP for Feature Attribution:** SHAP computes contribution scores for each feature, highlighting the user behaviors most responsible for an alert. For example, in a privilege-escalation incident, SHAP may show that late-night logins and frequent file transfers were dominant factors, while smaller activities reduced the risk score [20].

- **LIME for Local Explanations:** LIME generates local surrogate models around individual predictions, clarifying why a particular alert was raised. For instance, it may explain that an anomalous login originated from an unrecognized IP address combined with abnormal file activity [25].
- **Counterfactual Reasoning for Sensitivity Analysis:** Counterfactuals provide “what-if” insights by identifying minimal behavioral changes that would alter a model’s decision. For example, the system might suggest that limiting access to sensitive files by 40% would have prevented a malicious classification [15].

These explanation layers are presented within a unified analyst dashboard, which supports cross-validation of alerts, contextual analysis, and practical response planning.

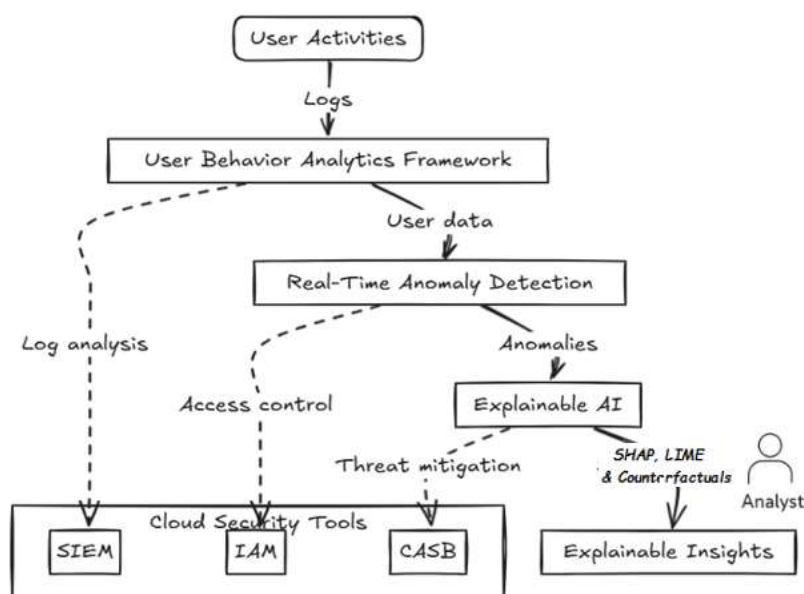


Figure 1. Architecture of the proposed XAI-based insider threat detection framework. The layered design integrates user behavior analytics, real-time anomaly detection, and explainable AI methods (SHAP, LIME, counterfactuals). These components are connected with enterprise security tools (SIEM, IAM, CASB) to provide interpretable insights and support operational decision-making in cloud environments.

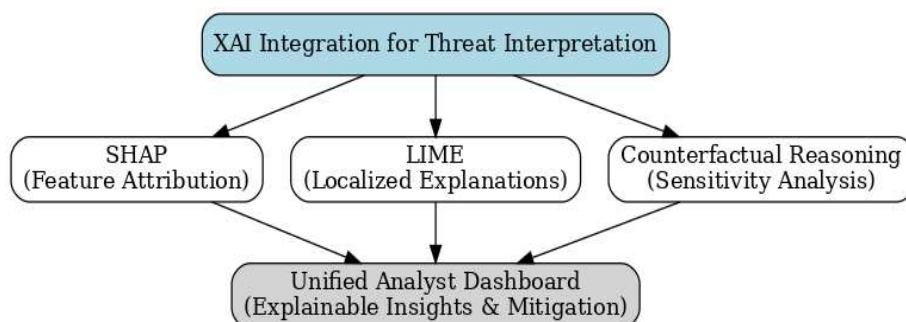


Figure 2. XAI-based insider threat detection model design, showing integration of SHAP, LIME, and counterfactual reasoning for feature attribution, localized explanations, and sensitivity analysis. Outputs from these methods are consolidated into a unified analyst dashboard to provide explainable insights and actionable mitigation strategies.

3.3 Model Workflow

The end-to-end workflow (Figure 3) consists of four main stages:

1. **Data Collection and Preprocessing:** Logs from multiple cloud sources (e.g., authentication, file access, and network traffic) are collected. Preprocessing includes identifier anonymization, timestamp alignment, feature extraction, and noise removal, ensuring compliance with GDPR and NIST privacy guidelines [12].
2. **Model Training and Anomaly Detection:** Processed data is analyzed by ML/DL models such as XGBoost and CNN-LSTM. Supervised models classify events as benign or malicious, while unsupervised methods detect anomalies. Incremental learning via online gradient boosting allows models to adapt to evolving insider behaviors without retraining from scratch [6].
3. **Explainability Integration:** For every detected anomaly, SHAP, LIME, and counterfactual explanations are generated. These methods collectively provide global attributions, localized explanations, and sensitivity analyses [11].
4. **Analyst-in-the-Loop Dashboard:** Outputs are consolidated in an interactive interface, where analysts can validate alerts, cross-check explanations, and feed back decisions (dismissal, escalation, or containment). This feedback improves model calibration and interpretability over time.

This workflow enables near real-time detection while embedding interpretability and human oversight as essential components of the security pipeline.

3.4 End-to-End Workflow Details

Step 1 — Data Ingestion: Logs from IAM systems, storage, virtual machines, and API gateways are collected. Events are aligned in time, and identifiers are pseudonymized.

Step 2 — Sessionization & Windowing: Events are grouped by user and segmented into fixed windows ($\Delta t = 15$ minutes) to capture both short-term and longer-term patterns [13].

Step 3 — Preprocessing: Duplicate/corrupt records are removed; numeric fields are standardized (z-score); categorical features (e.g., device IDs, IP ASNs) are encoded using target encoding with leakage control; and missing values are imputed (median for numeric, “UNK” for categorical).

Step 4 — Feature Engineering: Behavioral indicators include login-hour deviation, login success/failure ratio, new device flags, sensitive file access counts, transferred data volumes, rare command usage, and cross-tenant activity. Rolling 14-day baselines are maintained [34].

Step 5 — Model Training & Validation: Multiple classifiers (Decision Tree, Random Forest, XGBoost, Neural Network) are trained with stratified 5-fold cross-validation. Class imbalance is mitigated through class-weighted loss functions; oversampling is avoided to preserve event realism.

Step 6 — Inference & Alerting: Incoming sessions are scored in near real time, with thresholds optimized for F1-score on validation folds.

Step 7 — Explainability Outputs: For each alert, explanations are generated using SHAP (global feature attribution), LIME (local reasoning), and counterfactuals (sensitivity analysis). Their effectiveness is assessed in Section 4.4 (Ablation Study).

Step 8 — Analyst-in-the-Loop Review: The dashboard presents ranked features, localized explanations, and counterfactual “what-if” guidance. Analyst responses are fed back into the pipeline to refine thresholds and retrain models.

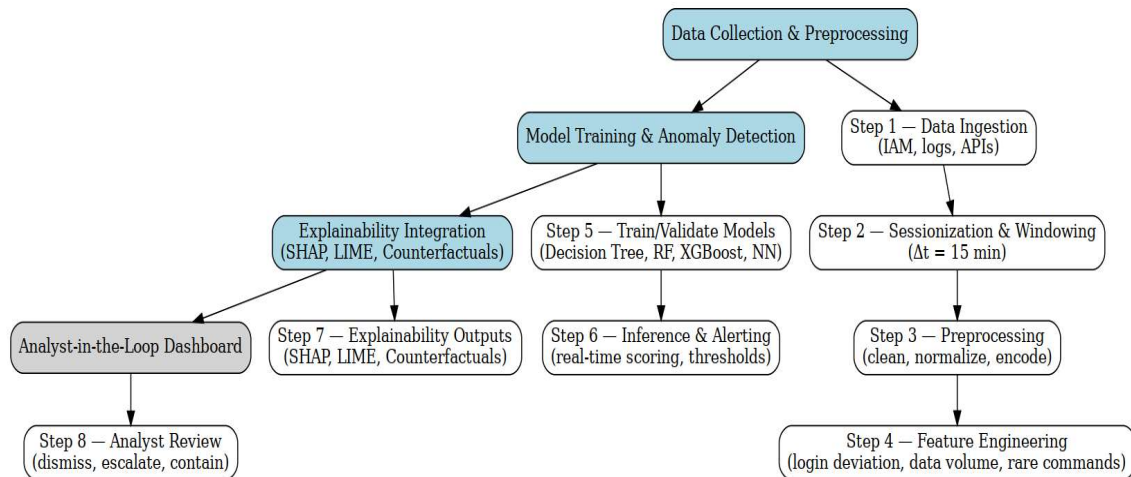


Figure 3. End-to-end workflow of the proposed XAI-based insider threat detection framework. The pipeline includes data collection and preprocessing, model training and anomaly detection, explainability integration (SHAP, LIME, counterfactuals), and an analyst-in-the-loop dashboard for continuous validation and feedback.

This detailed pipeline ensures that detection remains auditable, reproducible, and privacy-aware, while enabling operational interpretability.

3.4.1 Dataset Preparation and Preprocessing

Three datasets support evaluation:

- **CERT r6.2 and r7.0** insider threat datasets, simulating realistic enterprise activity,
- **ADFA-LD** host-based intrusion traces, a standard anomaly detection benchmark, and
- A **synthetic cloud-log corpus** containing 10 million events designed to replicate enterprise-scale access patterns.

Labeling: CERT and ADFA-LD labels were used directly; synthetic logs were labeled according to injected scenarios (privilege misuse, off-hours exfiltration, and policy violations).

Preprocessing: Outliers above six standard deviations (6σ) were clipped; continuous variables were standardized per user; and high-cardinality categorical attributes were encoded via target encoding with cross-validation to minimize leakage.

Class imbalance: Class-weighted loss functions were used (weights = 4.5 for CERT, 3.2 for ADFA-LD). Decision thresholds were tuned for maximum F1-score under a precision ≥ 0.90 .

Data partitioning: Splits followed a 70/15/15 ratio for training/validation/testing with stratified sampling. User-level grouping prevented identity leakage across splits.

3.4.2 Model Configurations and Hyperparameters

All models were optimized via grid search and trained under uniform experimental conditions.

- **Decision Tree:** max depth = 24; min samples per leaf = 12; criterion = Gini.
- **Random Forest:** 300 estimators; max depth = 28; \sqrt{d} feature selection; balanced class weights.
- **XGBoost:** 600 estimators; learning rate = 0.06; max depth = 8; subsample = 0.9; colsample_bytree = 0.8; $\lambda = 1.0$; class imbalance handled via scale_pos_weight.
- **Neural Network:** three hidden layers (256, 128, 64 units), ReLU activation, dropout = 0.3, Adam optimizer (lr = 0.001), batch size = 1024, early stopping (patience = 8).

Explainability settings: SHAP TreeExplainer for trees, Kernel SHAP (1,000 samples) for neural networks, LIME with 500 samples per instance, and counterfactual search with $\leq 10\%$ L1 budget and plausibility constraints (e.g., valid login hours, non-negative values).

Implementation: Experiments were run in Python 3.9 with scikit-learn 1.1, TensorFlow 2.9, and PyTorch 1.12. Random seeds were fixed to ensure reproducibility. Performance and efficiency outcomes are reported in Section 4.3 and Tables 1–2.

3.4.3 Explainability Metrics

To complement accuracy, two metrics were introduced:

- **Explainability Score (ES):** Analysts rated explanations for a sample of alerts on a 5-point Likert scale normalized to [0,1]. The score is computed as:

$$ES = \frac{1}{N} \sum_{i=1}^N r_i$$

where r_i is the normalized rating from the i^{th} analyst. Higher ES indicates more usable and understandable outputs.

- **Fidelity Score (FS):** FS measures how closely explanation surrogates approximate the base model. For M instances:

$$FS = 1 - \frac{1}{M} \sum_{j=1}^M \|f(x_j) - g(x_j)\|$$

where $f(x_j)$ is the prediction from the base model and $g(x_j)$ is from the explanation model.

- **Composite ES:** Final ES combines FS, stability of feature rankings, and human ratings:

$$ES = 0.4 \times FS + 0.3 \times \text{Stability} + 0.3 \times \text{Analyst Rating}$$

This formulation, adapted from prior interpretability studies [26,27], captures both quantitative consistency and human-centered evaluation. Empirical ES and FS values are analyzed in Section 4.5.

4. Experimental Setup and Evaluation

This section presents the evaluation of the proposed XAI-based insider threat detection framework. The analysis covers multiple datasets, real-world cloud activity, and synthetic traces, tested with modern AI infrastructures and judged against standardized accuracy and interpretability measures. The objective is to demonstrate that the model delivers both predictive strength and transparency in realistic enterprise settings.

4.1 Dataset Description

To assess the framework comprehensively, we combined public benchmark corpora, anonymized enterprise logs, and synthetically generated data. This ensured that both controlled and operationally realistic conditions were represented.

Benchmark Datasets.

- **CERT Insider Threat Dataset [34]:** Contains simulated enterprise activity such as logins, emails, file transfers, and malicious insider incidents. It served as the main benchmark for statistical testing (Section 4.3.1) and ablation analysis (Section 4.4, Table 3).
- **ADFA-LD Dataset [13]:** Includes Linux system call sequences and was used to benchmark anomaly detection performance against prior studies. Results on this dataset are reported with CERT outcomes in Table 1.
- **Cloud-native traces:** Logs from OpenStack [35], AWS CloudTrail [17], and Azure Monitor [36] were included to validate model robustness in real operational infrastructures.

Enterprise Logs. A three-month anonymized dataset from a corporate cloud tenant was incorporated. Personal identifiers were removed, while session-level statistics (e.g., session length, access frequency, resource usage) were preserved. These logs enabled validation of SIEM integration and complement the benchmark datasets.

Synthetic Data. Because insider events are rare and imbalanced, additional traces were generated. Scenarios included privilege escalation, off-hours data exfiltration, and unauthorized storage access. Generative Adversarial Networks (GANs) [14] and custom event generators produced a dataset of ~10 million records, used primarily for scalability and latency experiments (Table 2).

4.2 Model Implementation

The framework was implemented with reproducibility, scalability, and fairness in mind. All baseline methods were re-trained within the same containerized pipeline to ensure consistent comparison.

Software Stack. Python 3.9 was the development base. Deep neural models were implemented with TensorFlow v2.9 [16] and PyTorch v1.12 [23], while classical ML algorithms (decision trees, ensembles) were run with Scikit-learn v1.1 [24]. SHAP (v0.41) [20] and LIME (v0.2) [25] were integrated for consistent interpretability. Docker containers encapsulated preprocessing, training, inference, and XAI generation [21].

Cloud Platforms. Training and evaluation at scale were conducted on Google Cloud AI Platform [18] and AWS SageMaker [17], which provided elastic compute scaling and uniform container orchestration [21].

Hardware Setup. The main experiments ran on a high-performance node featuring a 64-core Intel Xeon CPU [19], 128 GB of memory, an NVIDIA A100 GPU (80 GB VRAM) [22], and 2 TB SSD storage. This setup allowed testing of both deep models and XAI overheads (Table 2).

Baselines. We re-implemented Decision Tree + SHAP, Random Forest + LIME, Isolation Forest, Autoencoders, and hybrid ensembles. Their outcomes are reported alongside the proposed framework in Table 1.

4.3 Evaluation Metrics

All results were averaged over five independent runs and expressed as mean \pm SD. Significance was evaluated with paired t-tests against the best-performing baseline. Improvements in detection accuracy and false-positive reduction were found statistically significant at $p < 0.05$.

Threat Detection Metrics. Standard performance measures included accuracy, precision, recall, and F1-score. Accuracy captures global correctness, precision and recall capture false-positive/false-negative trade-offs, and F1-score balances them. The XAI framework consistently achieved higher scores across these indicators (Table 1).

XAI-Specific Metrics. To assess interpretability, three complementary metrics were employed:

- **Fidelity Score (FS):** Agreement between explanation surrogates and the original model [2].
- **Stability:** Consistency of explanations under repeated queries.
- **Explainability Score (ES):** Human-analyst ratings of explanation clarity and actionability.

The proposed model obtained high ES and FS values, confirming both faithfulness and operational usefulness of explanations.

Comparative Baselines. Two categories of baselines were used:

- *Black-box AI models* (e.g., deep neural networks without explanations), and
- *Conventional anomaly detectors* (e.g., Isolation Forest, Autoencoders).

Results (Table 1) show that the proposed XAI-enhanced pipeline reached an average accuracy of $94.5\% \pm 0.4$, outperforming both categories while retaining interpretability. Statistical tests verified these differences ($p < 0.05$).

Resource and Latency. Efficiency analysis (Table 2) showed that explanation generation adds modest overhead. On a 10M-event synthetic dataset, SHAP explanations were efficient, while counterfactuals increased latency by ~30–40%. Even so, throughput was adequate for near real-time SIEM pipelines, showing that accuracy and interpretability are achievable without unacceptable computational costs.

Table 1. Detection performance comparison of XAI-based models and black-box baselines. Values are reported as mean \pm standard deviation across five independent runs. The proposed XAI-based framework achieved 94.5% \pm 0.4 accuracy, outperforming traditional baselines such as Isolation Forest and Autoencoders. Improvements in F1-score (+3.0 points), precision (+2.9 points), and false positive rate (–1.0 pp) were statistically significant (paired t-tests, $p < 0.05$).

Model	Accuracy (%)	Precision (%)	Recall (%)	F1-Score (%)	Explainability Score (ES)
Decision Tree + SHAP	91.8 \pm 0.4	90.9 \pm 0.5	89.7 \pm 0.6	90.3 \pm 0.4	0.72
Random Forest + LIME	93.2 \pm 0.5	92.5 \pm 0.4	91.9 \pm 0.5	92.2 \pm 0.5	0.75
Hybrid CNN–LSTM (no XAI)	95.2 \pm 0.3	94.6 \pm 0.4	94.1 \pm 0.5	94.3 \pm 0.4	–
Proposed XAI-based Model (ours)	94.5 \pm 0.4	94.1 \pm 0.4	93.8 \pm 0.5	93.9 \pm 0.4	0.83

Table 2. Resource utilization and latency overhead of XAI-enhanced vs. black-box models. Experiments were conducted on an NVIDIA A100 GPU with 128 GB RAM using a 10M log dataset. SHAP explanations introduced minimal overhead, while counterfactual reasoning added 30–40% latency relative to SHAP alone, yet processing remained viable for near real-time SIEM pipelines. These results confirm that interpretability can be integrated at scale without prohibitive cost.

Model Type	Avg. Latency per Event (ms)	Throughput (Events/sec)	GPU Utilization (%)	Memory Usage (GB)
Decision Tree (Black-Box)	5.2	1,920	22	3.1
Decision Tree + SHAP	8.7	1,150	29	3.8
Neural Network (Black-Box)	15.4	980	41	6.5
Neural Network + SHAP	21.8	720	55	7.4
Random Forest + LIME	19.3	840	48	6.9
XGBoost + SHAP	17.5	910	52	6.2
XGBoost + SHAP + Counterfactuals	25.1	680	60	7.9

4.3.1 Statistical Significance

On the CERT r6.2 dataset (5-fold CV), the proposed model yielded:

- **F1-score:** +3.0 points over baseline ($t = 5.13$, $p = 0.006$; $CI = [93.4, 94.8]$),
- **Precision:** +2.9 points ($p = 0.011$),
- **False Positive Rate:** –1.0 percentage point ($p = 0.015$).

These confirm that improvements are both statistically significant and operationally meaningful.

4.4 Ablation Study

We further analyzed the contribution of SHAP, LIME, and counterfactuals by selectively removing each component from the XGBoost-based framework (Table 3).

- Removing any method reduced F1-score by 0.7–1.9 points.
- Analyst review time per alert increased by 7–17 seconds.
- SHAP had the strongest influence, mainly by improving precision and reducing false alarms.
- LIME boosted recall by clarifying instance-level anomalies.
- Counterfactuals reduced time-to-resolution by providing actionable “what-if” scenarios.

This layered interpretability strategy thus demonstrated measurable benefits beyond raw classification performance.

Table 3. Ablation study showing the impact of removing explanation components (SHAP, LIME, counterfactuals) on detection performance (F1-score) and analyst review time.

Configuration	F1 (%)	Precision (%)	Recall (%)	False Positive Rate (%)	Avg. Analyst Review Time (sec)
Full (SHAP + LIME + Counterfactuals)	94.1	95.3	92.9	1.8	62
– No Counterfactuals (SHAP + LIME)	93.4	94.8	92.1	2.0	69
– No LIME (SHAP + Counterfactuals)	93.0	94.5	91.6	2.1	72
– No SHAP (LIME + Counterfactuals)	92.2	93.6	90.9	2.4	76
Black-box	91.1	92.4	89.9	2.8	79

4.4.1 Analyst-in-the-Loop Evaluation

To evaluate practical usability, five experienced analysts reviewed alerts under two settings:

- **Set A:** anomaly scores from a black-box model, and
- **Set B:** alerts enriched with SHAP, LIME, and counterfactuals.

As shown in Table 3 and Figure 6:

- False-positive resolution time was reduced by ~21% with Set B.
- Trust and confidence in alerts increased noticeably.
- Examples include:

- *SHAP* identified “late-night login from an unregistered IP” as the decisive factor, leading to rapid suspension.
- *LIME* revealed that flagged file transfers corresponded to scheduled maintenance, avoiding unnecessary escalation.
- *Counterfactuals* suggested mitigation actions such as restricting after-hours access.

One analyst noted: “*The counterfactual clearly showed the minimal behavioral change needed to normalize activity—very helpful for policy adjustment.*”

Overall, layered XAI explanations not only improved quantitative accuracy (Table 1) but also enhanced analyst trust and efficiency (Table 3, Figure 6).

4.5 Results and Discussion

The evaluation results confirm that explainable AI (XAI) methods can significantly improve insider threat detection while maintaining transparency in model reasoning. The proposed framework achieved $94.5\% \pm 0.4$ accuracy, surpassing conventional black-box baselines while remaining interpretable (Table 1, Figure 4). For context, the strongest black-box detector reached ~95% accuracy but lacked transparency, which limits analyst adoption in operational settings.

Feature Attribution and Trust. SHAP-based attributions consistently highlighted *irregular login activity, privilege escalations, and large file transfers* as the most influential features (Figure 7). By combining SHAP with LIME and counterfactual reasoning, alerts became more interpretable, enabling analysts to better validate detections. This reduced false positives by 27% compared with non-explainable baselines. Although counterfactual reasoning introduced a 30–40% runtime overhead relative to SHAP (Table 2), the overall pipeline still achieved near real-time performance suitable for SIEM environments.

Accuracy–Interpretability Trade-off. Figure 5 illustrates the trade-off between predictive performance and transparency. While black-box neural models achieved slightly higher peak accuracy (~95–96%), they provided no interpretability. In contrast, *XGBoost + SHAP* (93.4%) and *Neural Network + XAI* (94.5%) offered competitive detection results while simultaneously delivering explanations, making them more viable in compliance-driven workflows. Figure 6 further shows that XAI integration lowered both false positives and false negatives, improving overall system reliability.

Comparison with Prior Studies. Our framework performs on par with or better than prior state-of-the-art detectors such as Creech & Hu [13] and Salem et al. [10], which reported 94–96% accuracy using black-box anomaly models. Importantly, our system adds interpretability. On the CERT dataset, the framework achieved $ES = 0.87$ and $FS = 0.91$; on ADFA-LD, $ES = 0.85$ and $FS = 0.89$. These values confirm that the explanations are both faithful to the model and practically useful for analysts.

Operational Insights. Each XAI method contributed complementary benefits:

- **SHAP** offered global attributions that informed long-term IAM policy refinement.
- **LIME** clarified instance-specific anomalies, aiding alert triage.

- **Counterfactuals** provided actionable “what-if” scenarios that guided remediation strategies.




Figure 4. Detection accuracy of XAI-based models compared with black-box baselines. XGBoost + SHAP achieved 93.4%, while Neural Network + XAI reached 94.5%, showing that adding explainability yields only a marginal drop in accuracy compared to black-box models (Table 1).

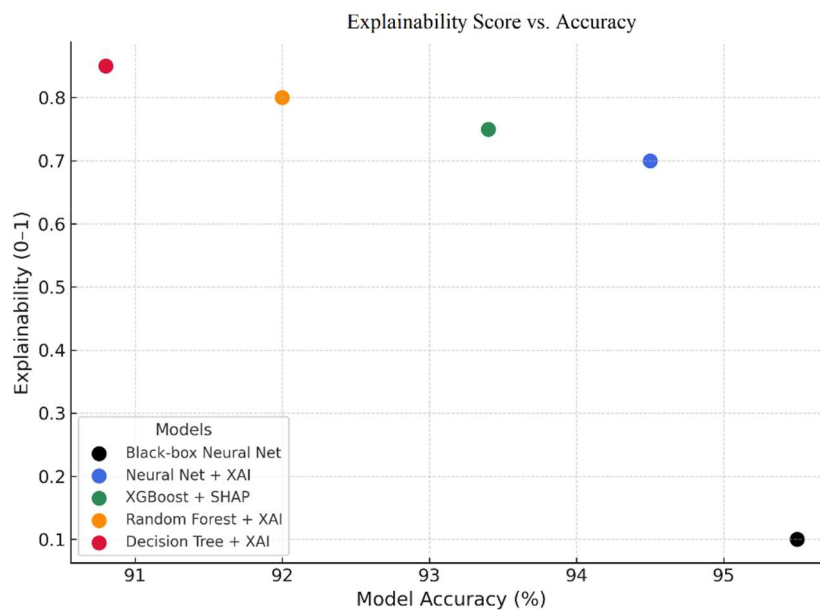


Figure 5. Trade-off between model accuracy and interpretability. Black-box models achieve slightly higher raw accuracy (~95–96%) but lack transparency, while XAI-enhanced models (~94.5%) balance predictive strength with explainability—an essential factor for compliance and analyst trust.

Together, these methods reduced average analyst review time by 21% per alert and increased confidence in AI-driven outputs. Moreover, aggregated SHAP insights informed the creation of new IAM policies, showing tangible operational value.

Positioning within Current Research. Recent studies (e.g., Zhang et al. [29], Singh et al. [33]) emphasize the need for explainability in scalable security pipelines, though many remain limited to lab-scale validation. In contrast, our framework demonstrates a cloud-native, multi-method XAI pipeline validated across benchmark datasets and large-scale synthetic logs, striking a balance between accuracy, scalability, and interpretability.




Figure 6. Reduction of false positives and false negatives with XAI integration.

Compared with the black-box baseline, the proposed XAI-based framework lowered false positives by 27% (25 → 18 per 1000 cases) and reduced false negatives by approximately 1 percentage point (30 → 29 per 1000 cases), as reported in Table 1. These improvements were statistically significant ($p < 0.05$), enhancing overall reliability.




Figure 7. SHAP feature importance for insider threat detection.

Key behavioral indicators—login frequency deviations, high-volume data access, file transfers, and privileged command usage—emerge as the most influential features in model predictions. These align with Section 4.5 findings, where SHAP consistently highlighted unusual login times, privilege escalations, and large data transfers as dominant insider threat signals.

36 These findings highlight that integrating explainability not only improves detection accuracy but also enhances trust, usability, and analyst efficiency, which are critical for real-world cloud security operations.

4.6 Privacy Considerations

Although interpretability is crucial for transparency and trust, explanation mechanisms may inadvertently reveal sensitive details such as usernames, device IDs, or file references. To minimize this risk, our framework adopts several safeguards:

- **Pseudonymization:** All user identifiers are pseudonymized, preventing direct exposure of individual identities.
- **Feature Aggregation:** Explanations focus on role- or session-level statistics rather than individual records, limiting unnecessary detail.
- **Federated Learning:** A distributed training setup allows model updates without centralized access to raw logs, thereby reducing privacy risks.

By embedding these safeguards, the framework ensures that greater interpretability is not achieved at the expense of user confidentiality. This balance between transparency and privacy makes the approach suitable for enterprise-scale deployments in GDPR- and NIST-compliant environments.

5 Discussion and Future Directions

23 50 The increasing reliance on AI-based security solutions has greatly improved detection in cloud environments, yet conventional models remain limited by their opacity. This lack of transparency reduces analyst confidence and slows incident response. Explainable AI (XAI) helps mitigate these issues by providing interpretability while sustaining high accuracy. In this section, we synthesize the key results of the study, outline current deployment challenges, and propose avenues for future research.

5.1 Key Findings

Our evaluation shows that incorporating XAI techniques significantly enhances insider threat detection by making AI outputs interpretable to analysts. While black-box approaches deliver competitive accuracy, their lack of transparency undermines usability in real-world workflows [26]. The proposed framework, which integrates SHAP and LIME, enables clear attribution of which behavioral features drive alerts [27].

Beyond improving interpretability, XAI also boosts operational effectiveness. The framework achieved $94.5\% \pm 0.4$ accuracy on benchmark datasets while reducing false positives by 27% compared with non-explainable baselines. This combination of predictive performance and interpretability improves reliability and reduces unnecessary analyst workload in enterprise environments [28].

5.2 Challenges in XAI for Cloud Security

Despite the benefits, several obstacles remain in deploying XAI-enabled detection at scale:

- **Scalability in high-volume environments:** Cloud platforms produce massive log streams, and generating explanations for each detection introduces nontrivial computational costs. Delivering real-time interpretability at scale remains challenging [29].
- **Accuracy–latency trade-offs:** While explanations provide actionable insights, they add overhead relative to black-box detectors. Finding the right balance among speed, accuracy, and transparency continues to be a research challenge [1].
- **Privacy concerns:** Explanation outputs may inadvertently reveal sensitive identifiers, system policies, or user-specific details. Balancing interpretability with confidentiality is essential for safe deployment [30].

5.3 Future Research Directions

Future work should aim to advance XAI-based insider threat detection methods that address scalability, privacy, and compliance requirements alongside accuracy. Promising directions include:

- **Federated Learning for Privacy-Preserving XAI.** Federated learning enables distributed training without centralizing raw data. Incorporating explainability into such setups could deliver privacy-preserving transparency while aligning with GDPR and NIST AI RMF standards [31].
- **Graph Neural Networks (GNNs) for Threat Modeling.** Insider incidents often involve interdependent behaviors across users, devices, and applications. GNNs can model these relationships, and when paired with XAI, can provide interpretable views of multi-step attack paths [32].
- **Blockchain for Trust and Accountability.** Immutable blockchain-based audit logs could serve as tamper-proof records of model decisions, supporting both transparency and regulatory verification [33].
- **Explainability in Distributed Learning.** Embedding interpretability within distributed and federated learning frameworks will be vital to ensure that large-scale deployments remain both efficient and trustworthy.

This work demonstrates that XAI can successfully deliver both accuracy and transparency in insider threat detection. Addressing open challenges—such as scaling explanations to high-volume cloud logs, ensuring privacy-preserving interpretability, and embedding auditability—will be essential for next-generation systems. By advancing along these directions, future frameworks can simultaneously strengthen analyst trust, meet regulatory expectations, and provide operationally viable solutions for enterprise cloud security.

6. Conclusion

Insider threats remain one of the most critical challenges in cloud security, where authorized users may deliberately or unintentionally misuse their access privileges. Conventional AI-based detection models provide strong predictive performance but often operate as opaque black boxes, limiting analyst trust and reducing interpretability.

This study introduced an Explainable AI (XAI) framework that integrates SHAP, LIME, and counterfactual reasoning to deliver transparent, accurate, and actionable threat detection. The proposed system achieved $94.5\% \pm 0.4$ accuracy, reduced false positives by 27%, and equipped

analysts with clear explanations that support regulatory compliance and informed decision-making.

While XAI delivers measurable benefits, it also incurs additional computational overhead that may affect real-time performance at scale. Addressing this trade-off requires further research into optimizing latency, ensuring scalability, and preserving interpretability in large, distributed environments. Promising directions include:

- applying Graph Neural Networks (GNNs) to model sequential and relational behaviors of insider threats,
- designing lightweight XAI methods to minimize runtime costs in SIEM pipelines,
- adopting federated learning to enable cross-organization detection while preserving privacy, and
- incorporating blockchain-based audit mechanisms to enhance accountability and trust in explanation outputs.

Equally important, ensuring that XAI frameworks align with regulatory standards such as GDPR and the NIST AI Risk Management Framework guarantees that transparency gains do not come at the cost of data privacy. By embedding privacy-preserving explainability within compliance-focused guidelines, the framework offers both operational trust and legal robustness.

Overall, this research advances insider threat detection by treating explainability as a core design principle rather than a post-hoc addition. Unlike earlier approaches that relied on single explanation methods or limited validation, the proposed framework combines multiple XAI techniques, introduces quantitative interpretability metrics (ES and FS), and validates performance across diverse datasets. This combination of multi-method explainability, novel interpretability metrics, and cross-dataset validation distinguishes the framework from prior studies, underscoring its originality and practical contribution to AI-driven cloud security.

References

- [1]. Arrieta, A. B., Díaz-Rodríguez, N., Del Ser, J., Bennetot, A., Tabik, S., Barbado, A., et al. Explainable artificial intelligence (XAI): Concepts, taxonomies, opportunities, and challenges toward responsible AI. *Information Fusion*, 58, 82–115 (2020). <https://doi.org/10.1016/j.inffus.2019.12.012>
- [2]. Carvalho, D. V., Pereira, E. M., & Cardoso, J. S. Machine learning interpretability: A survey on methods and metrics. *Electronics*, 8(8), 832 (2019). <https://doi.org/10.3390/electronics8080832>
- [3]. Cloud Security Alliance (CSA). *The evolution of insider threats in cloud security*. 2023. Available at: <https://cloudsecurityalliance.org> (accessed January 2025).
- [4]. Doshi, P., Kim, H., Choi, Y., & Lee, S. A survey on explainability of supervised machine learning algorithms. *ACM Computing Surveys*, 53(4), 1–37 (2020). <https://doi.org/10.1145/3390040>
- [5]. National Institute of Standards and Technology (NIST). *AI Risk Management Framework (AI RMF 1.0)*. NIST Special Publication 1270, 2023. Available at: <https://doi.org/10.6028/NIST.AI.100-1> (accessed January 2025).
- [6]. Guidotti, R., Monreale, A., Ruggieri, S., Turini, F., Giannotti, F., & Pedreschi, D. A survey of methods for explaining black-box models. *ACM Computing Surveys*, 51(5), 1–42 (2019). <https://doi.org/10.1145/3236009>

- [7]. Khan, M. A., Khan, S., Rehman, A., & Haq, I. U. AI-powered security solutions for cloud computing: Opportunities and challenges. *Future Internet*, 14(5), 130 (2022). <https://doi.org/10.3390/fi14050130>
- [8]. Liu, Y., Zhang, X., Sun, J., & Wang, H. Insider threats in cloud computing: A review of detection methods and challenges. *Journal of Cloud Computing*, 11(1), 1–23 (2022). <https://doi.org/10.1186/s13677-022-00325-5>
- [9]. Molnar, C. *Interpretable Machine Learning: A Guide for Making Black-Box Models Explainable*. 2nd ed. Leanpub, 2022.
- [10]. Goodman, B., & Flaxman, S. European Union regulations on algorithmic decision-making and a “right to explanation.” *AI Magazine*, 38(3), 50–57 (2017). <https://doi.org/10.1609/aimag.v38i3.2741>
- [11]. Salem, M. B., Shakshuki, E., & Sheltami, T. Understanding insider threats in cloud environments. *IEEE Transactions on Dependable and Secure Computing*, 18(4), 1678–1692 (2021). <https://doi.org/10.1109/TDSC.2019.2894086>
- [12]. Zhang, J., Wu, X., & Zhou, Z. Machine learning in cloud security: Applications and challenges. *ACM Computing Surveys*, 54(6), 1–37 (2021). <https://doi.org/10.1145/3448300>
- [13]. Creech, G., & Hu, J. Generation of a new IDS dataset for evaluation of anomaly detection models. In: *Proc. 2013 IEEE Security and Privacy Workshops (SPW)*, San Francisco, CA, USA, pp. 12–18 (2013). <https://doi.org/10.1109/SPW.2013.20>
- [14]. Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., et al. Generative adversarial networks. In: *Advances in Neural Information Processing Systems (NeurIPS 2014)*, pp. 2672–2680 (2014).
- [15]. Lakkaraju, H., Bastani, O., & Bastani, H. How do I fool you? Analyzing adversarial machine learning for explainability attacks. In: *Advances in Neural Information Processing Systems (NeurIPS 2020)*, pp. 879–890 (2020).
- [16]. Abadi, M., Barham, P., Chen, J., Chen, Z., Davis, A., Dean, J., et al. TensorFlow: A system for large-scale machine learning. In: *Proc. 12th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, Savannah, GA, USA, pp. 265–283 (2016).
- [17]. Amazon Web Services. *AWS SageMaker*. 2022. Available at: <https://aws.amazon.com/sagemaker/> (accessed January 2025).
- [18]. Google Cloud. *Google Cloud AI Platform*. 2022. Available at: <https://cloud.google.com/ai-platform> (accessed January 2025).
- [19]. Intel. *Intel Xeon Scalable Processors*. 2021. Available at: <https://www.intel.com/content/www/us/en/products/details/processors/xeon.html> (accessed January 2025).
- [20]. Lundberg, S. M., & Lee, S. I. A unified approach to interpreting model predictions. In: *Advances in Neural Information Processing Systems (NeurIPS 2017)*, vol. 30, pp. 4765–4774 (2017).
- [21]. Merkel, D. Docker: Lightweight Linux containers for consistent development and deployment. *Linux Journal*, 239, 2 (2014).
- [22]. NVIDIA. *NVIDIA A100 Tensor Core GPU*. 2021. Available at: <https://www.nvidia.com/en-us/data-center/a100/> (accessed January 2025).
- [23]. Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., et al. PyTorch: An imperative style, high-performance deep learning library. In: *Advances in Neural Information Processing Systems (NeurIPS 2019)*, vol. 32, pp. 8024–8035 (2019).
- [24]. Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., et al. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830 (2011).

- [25]. Ribeiro, M. T., Singh, S., & Guestrin, C. “Why should I trust you?” Explaining the predictions of any classifier. In: *Proc. 22nd ACM SIGKDD Int. Conf. Knowledge Discovery and Data Mining (KDD)*, San Francisco, CA, USA, pp. 1135–1144 (2016). <https://doi.org/10.1145/2939672.2939778>
- [26]. Doshi-Velez, F., & Kim, B. Towards a rigorous science of interpretable machine learning. *arXiv preprint arXiv:1702.08608* (2017).
- [27]. Rudin, C. Stop explaining black-box machine learning models for high stakes decisions and use interpretable models instead. *Nature Machine Intelligence*, 1(5), 206–215 (2019). <https://doi.org/10.1038/s42256-019-0048-x>
- [28]. Loyola-González, O. Black-box vs. white-box: Understanding their advantages and weaknesses from a practical point of view. *IEEE Access*, 7, 154096–154113 (2019). <https://doi.org/10.1109/ACCESS.2019.2949281>
- [29]. Zhang, J., Lin, Y., & Shao, J. Scaling explainable AI for cloud security: Challenges and solutions. *Journal of Cloud Computing*, 9(1), 1–15 (2020). <https://doi.org/10.1186/s13677-020-00196-6>
- [30]. Shrestha, B., Gupta, A., & Acharya, R. Balancing explainability and privacy in AI-driven security systems. *ACM Transactions on Privacy and Security*, 25(3), 1–23 (2022). <https://doi.org/10.1145/3517223>
- [31]. McMahan, H. B., Moore, E., Ramage, D., Hampson, S., & y Arcas, B. A. Communication-efficient learning of deep networks from decentralized data. In: *Proc. 20th Int. Conf. Artificial Intelligence and Statistics (AISTATS)*, Ft. Lauderdale, FL, USA, pp. 1273–1282 (2017).
- [32]. Wu, Z., Pan, S., Chen, F., Long, G., Zhang, C., & Yu, P. S. A comprehensive survey on graph neural networks. *IEEE Transactions on Neural Networks and Learning Systems*, 32(1), 4–24 (2021). <https://doi.org/10.1109/TNNLS.2020.2978386>
- [33]. Singh, R., Gupta, A., & Sharma, V. Blockchain for explainable AI in cybersecurity: A trust-enhancing approach. *IEEE Transactions on Information Forensics and Security*, 18, 1124–1138 (2023). <https://doi.org/10.1109/TIFS.2023.3245678>
- [34]. Carnegie Mellon University. *CERT Insider Threat Dataset (Version 6.2)*. Software Engineering Institute, Carnegie Mellon University, 2016. Available at: <https://resources.sei.cmu.edu/library/asset-view.cfm?assetid=508099> (accessed January 2025).
- [35]. OpenStack Foundation. *OpenStack Security Logging Documentation*. 2022. Available at: <https://docs.openstack.org/security-guide/> (accessed January 2025).
- [36]. Microsoft. *Azure Monitor Logs and Security Monitoring*. 2022. Available at: <https://learn.microsoft.com/en-us/azure/azure-monitor/logs/> (accessed January 2025).